

Fraud Detection prediction using XGBoosting

Haijie CAO

10/05/2017

Reading Data

```
card<-read.csv("creditcard.csv")

#card$Class<-as.factor(card$Class)
table(card$Class)

##
##      0      1
## 284315  492

taux<-nrow(card[card$Class==1,])/nrow(card)
taux

## [1] 0.001727486

## separete the data into training set and testing set
cardsub<-card[,-1]
set.seed(5)
ind<-sample(2,nrow(cardsub),replace = T,prob=c(0.7,0.3))
train<-cardsub[ind==1,]
test<-cardsub[ind==2,]

#####
feature.names=names(train)

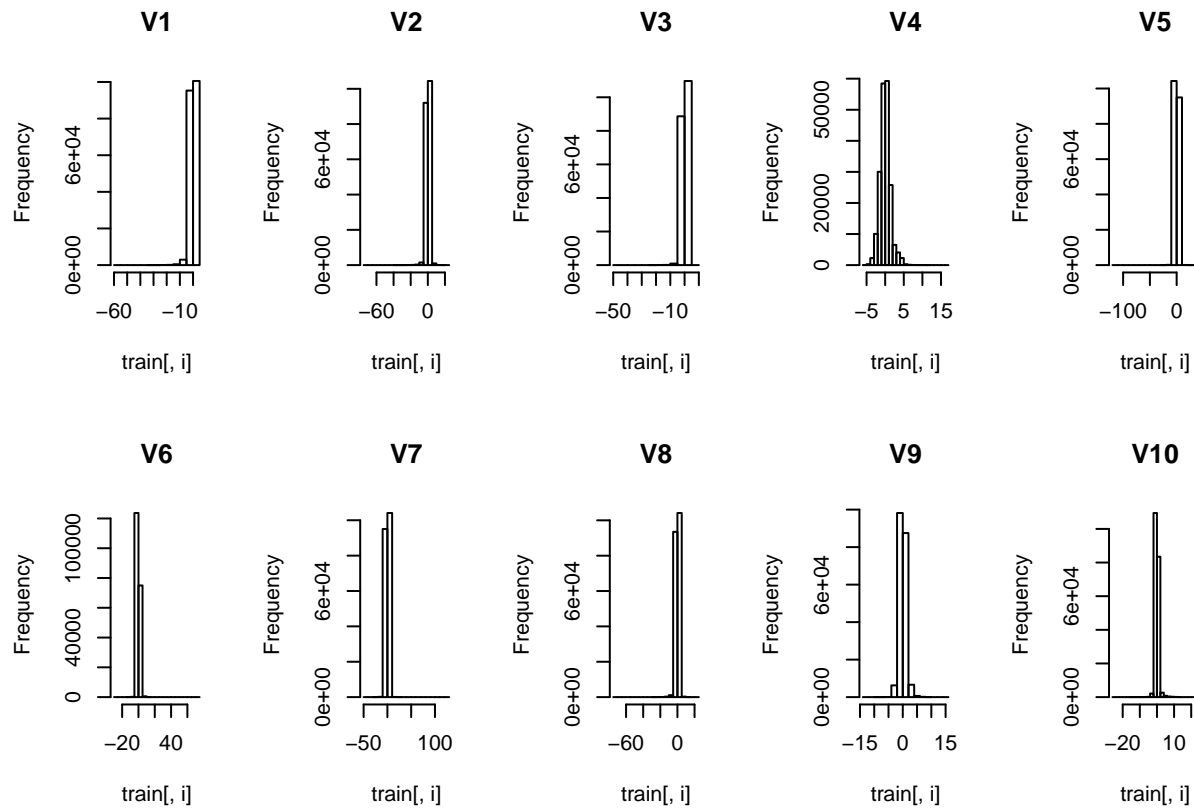
for (f in feature.names) {
  if (class(train[[f]])=="factor") {
    levels <- unique(c(train[[f]]))
    train[[f]] <- factor(train[[f]],
                        labels=make.names(levels))
  }
}

for (f in feature.names) {
  if (class(test[[f]])=="factor") {
    levels <- unique(c(test[[f]]))
    test[[f]] <- factor(test[[f]],
                      labels=make.names(levels))
  }
}
#####
```

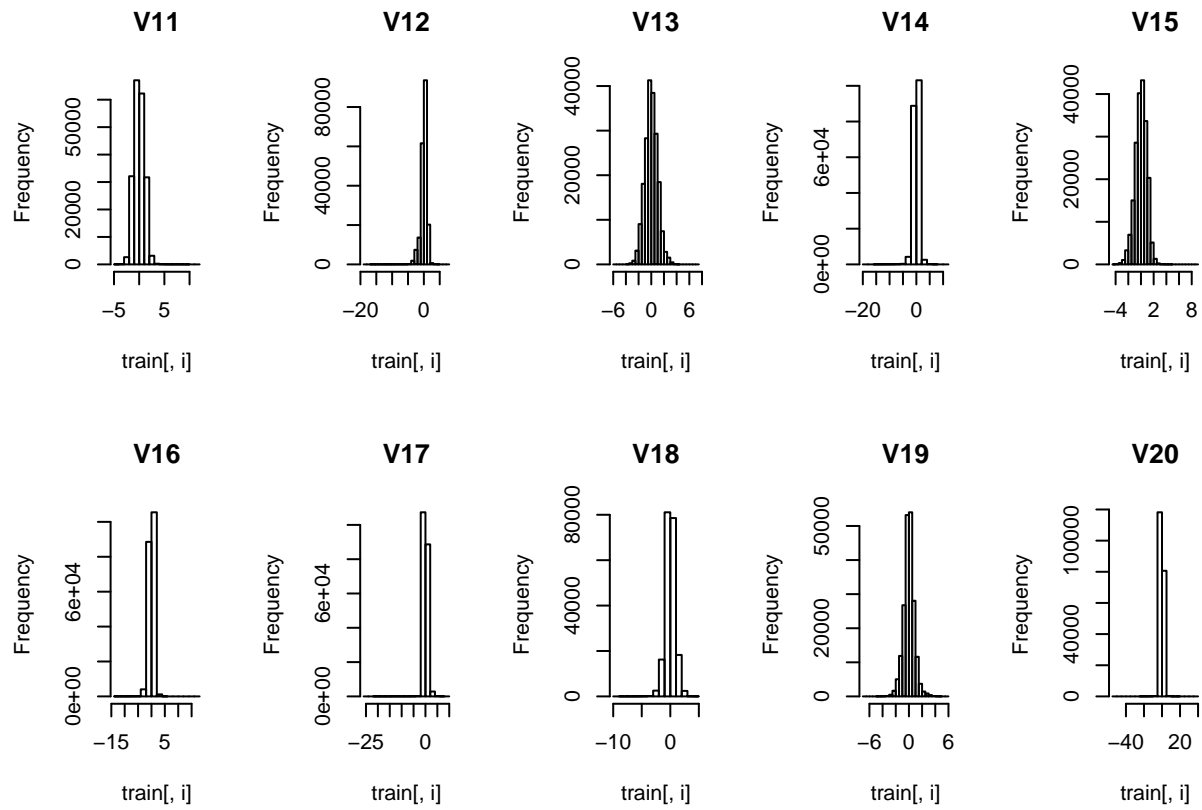
Dada Exploratory

```
par(mfrow=c(2,5))
for(i in 1:10){
```

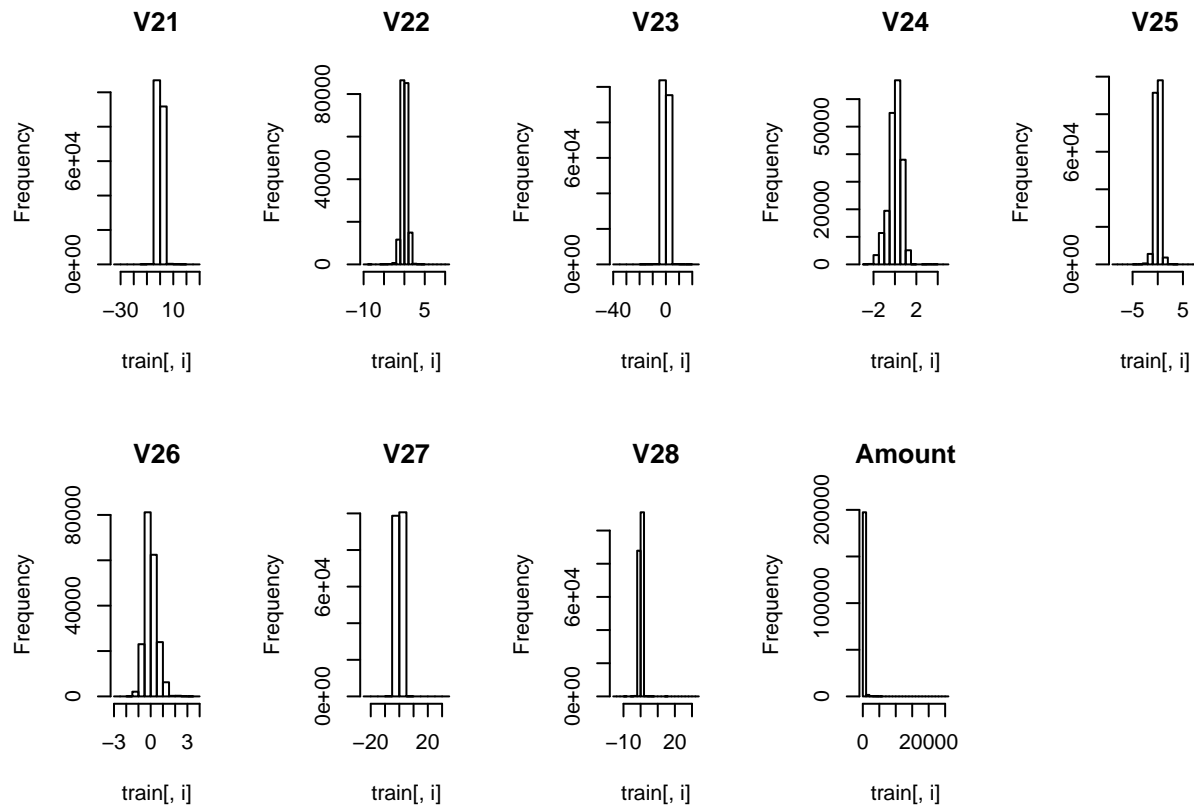
```
hist(train[,i], main=names(train)[i])
}
```



```
par(mfrow=c(2,5))
for(i in 11:20){
  hist(train[,i], main=names(train)[i])
}
```



```
par(mfrow=c(2,5))
for(i in 21:29){
  hist(train[,i], main=names(train)[i])
}
```



We can see that the variables are very skewed.

XGBoost

```
library(ggplot2,quietly = T)
library(plyr)
library(dplyr,quietly = T)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:plyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(tidyr,quietly = T)
library(readr,quietly = T)
library(xgboost,quietly = T)

##
```

```

## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##      slice

library(caret,quietly = T)
library(pROC,quietly = T)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

# xgboost fitting with arbitrary parameters
xgb_params = list(
  objective = "binary:logistic",      # binary classification
  eta = 0.01,                        # learning rate
  max.depth = 3,                     # max tree depth
  eval_metric = "auc"                # evaluation/loss metric
)

# fit the model with the arbitrary parameters specified above
xgb = xgboost(data = as.matrix(train[,1:29]),
  label = train$Class,
  params = xgb_params,
  nrounds = 100,                      # max number of trees to build
  verbose = TRUE,
  print_every_n = 20,
  early_stop_round = 10              # stop if no improvement within 10 trees
)

## [1] train-auc:0.909587
## [21] train-auc:0.912663
## [41] train-auc:0.912664
## [61] train-auc:0.912673
## [81] train-auc:0.912677
## [100] train-auc:0.912698

# cross-validate xgboost to get the accurate measure of error
xgb_cv = xgb.cv(params = xgb_params,
  data = as.matrix(train[,1:29]),
  label = train$Class,
  nrounds = 100,
  nfold = 5,                          # number of folds in K-fold
  prediction = TRUE,                  # return the prediction using the final model
  showsd = TRUE,                     # standard deviation of loss across folds
  stratified = TRUE,                  # sample is unbalanced; use stratified sampling
  verbose = TRUE,
  print_every_n = 20,
  early_stop_round = 10
)

## [1] train-auc:0.909600+0.003094 test-auc:0.904845+0.007888

```

```
## [21] train-auc:0.911566+0.002337 test-auc:0.904872+0.007886
## [41] train-auc:0.912704+0.003098 test-auc:0.909549+0.006195
## [61] train-auc:0.913846+0.003201 test-auc:0.911158+0.008934
## [81] train-auc:0.913859+0.003198 test-auc:0.911169+0.008936
## [100] train-auc:0.914230+0.002974 test-auc:0.911171+0.008934
```

```
# set up the cross-validated hyper-parameter search
```

```
xgb_grid = expand.grid(
  nrounds = 100,
  eta = c(0.1),
  max_depth = c(2,6),
  gamma = 1,
  colsample_bytree=1,
  min_child_weight=10,
  subsample=1)
```

```
# pack the training control parameters
```

```
xgb_trcontrol = trainControl(
  method = "cv",
  number = 5,
  verboseIter = TRUE,
  returnData = FALSE,
  returnResamp = "all",
  classProbs = TRUE,
  summaryFunction = twoClassSummary,
  allowParallel = TRUE
```

```
# save losses across all models
```

```
# set to TRUE for AUC to be computed
```

```
)
```

```
# train the model for each parameter combination in the grid,
```

```
# using CV to evaluate
```

```
train$Class[train$Class==0]<-"No"
```

```
train$Class[train$Class==1]<-"Yes"
```


```
xgb_train = train(
  x = as.matrix(train[,1:29]),
  y = as.factor(train$Class),
  trControl = xgb_trcontrol,
  metric="ROC",
  tuneGrid = xgb_grid,
  method = "xgbTree"
```

```
)
```

```
## + Fold1: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold1: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold1: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold1: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold2: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold2: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold2: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold2: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold3: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold3: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold3: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold3: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
```

```
## + Fold4: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold4: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold4: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold4: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold5: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold5: eta=0.1, max_depth=2, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## + Fold5: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## - Fold5: eta=0.1, max_depth=6, gamma=1, colsample_bytree=1, min_child_weight=10, subsample=1, nrounds=100
## Aggregating results
## Selecting tuning parameters
## Fitting nrounds = 100, max_depth = 6, eta = 0.1, gamma = 1, colsample_bytree = 1, min_child_weight = 10
xgb_train$bestTune
```

```
##   nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 2      100      6 0.1    1          1          10          1
plot(xgb_train)
```



```
res <- xgb_train$results
res
```

```
##   eta max_depth gamma colsample_bytree min_child_weight subsample nrounds
## 1 0.1         2     1           1           10           1      100
## 2 0.1         6     1           1           10           1      100
##           ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1 0.9753148 0.9998042 0.7430303 0.01280935 7.615431e-05 0.09711998
## 2 0.9762149 0.9998644 0.7676457 0.01154988 2.862677e-05 0.08889640
```

```
### xgboostModel Predictions and Performance
# Make predictions using the test data set
xgb.pred <- predict(xgb_train,test)
```

```

test$Class[test$Class==0]<-"No"
test$Class[test$Class==1]<-"Yes"
test$Class<-as.factor(test$Class)

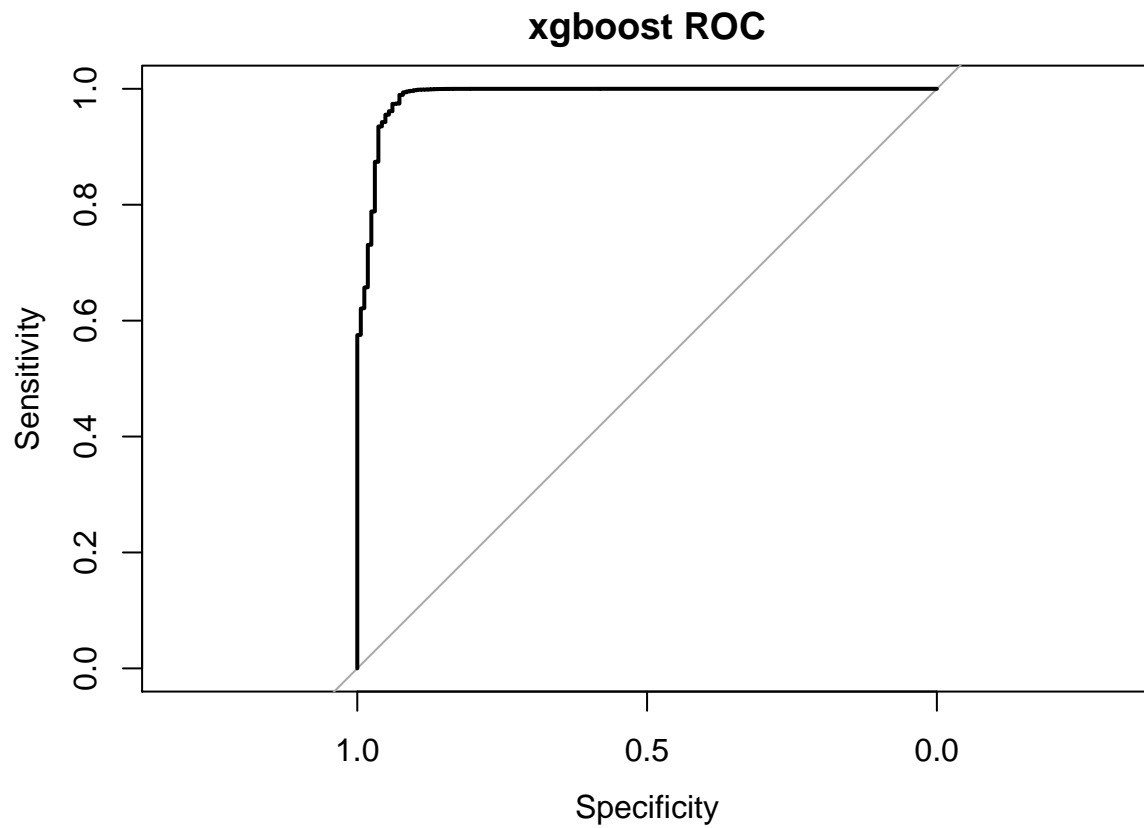
#Look at the confusion matrix
confusionMatrix(xgb.pred,test$Class)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    No    Yes
##           No 85164    29
##           Yes   6   136
##
##           Accuracy : 0.9996
##           95% CI : (0.9994, 0.9997)
##           No Information Rate : 0.9981
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8858
##           Mcnemar's Test P-Value : 0.0002003
##
##           Sensitivity : 0.9999
##           Specificity : 0.8242
##           Pos Pred Value : 0.9997
##           Neg Pred Value : 0.9577
##           Prevalence : 0.9981
##           Detection Rate : 0.9980
##           Detection Prevalence : 0.9983
##           Balanced Accuracy : 0.9121
##
##           'Positive' Class : No
##
#Draw the ROC curve
xgb.probs <- predict(xgb_train,test,type="prob")

#head(xgb.probs)
xgb.ROC <- roc(predictor=xgb.probs$No,
               response=test$Class,
               levels=rev(levels(test$Class)))
xgb.ROC$auc

## Area under the curve: 0.9876
plot(xgb.ROC,main="xgboost ROC")

```

```
# Plot the propability of poor segmentation
histogram(~xgb.probs$No|test$Class,xlab="Probability of Poor Segmentation")
```

